


# Scalable System

Common Problems

**Asmit Ghimire**  
Software Engineer at Gurzu Inc.

# What does Scalability mean for Systems and Services?

- ❖ System that can handle rapid changes to workloads and user demands.



# Premature Scaling

- ❖ Premature scaling happens when startups try to scale up too early.
- ❖ It can kill a startup company with unplanned budgeting.

Proper steps to scale:

1. Discover
2. Validate
3. Scale



# Problems Faced in a Scaling System

❖ Issues encountered as of my experience:


- Insufficient system infrastructure
- N+1 queries
- Memory leak / bloat
- Bulky / slow APIs
- Missing DB indexes

# Server Scaling

- ❖ There are two ways we can go about scaling servers:
  - Vertical Scaling
  - Horizontal Scaling


## Vertical Scaling

( Increase size of instance (RAM , CPU etc.) )



## Horizontal Scaling

( Add more instances )



# Vertical Scaling

- ❖ Increase the capability of a single server.
- ❖ Appropriate for a growing system.


## Pros:

- Easy to implement and maintain
- Cost effective

## Cons:

- Single point of failure
- Hardware limitations

( Increase size of instance (RAM , CPU etc.) )



# Horizontal Scaling

- ❖ Add more of the same instance of resources.
- ❖ Workload distributed by load balancers.
- ❖ Appropriate for a large system.


Pros:

- More fault tolerance and fewer risks of downtime
- Scales well with increasing demand

Cons:


- Load balancing is required
- Expensive

( Add more instances )



# Load Balancing

- ❖ LB distributes incoming network traffic across multiple servers.
- ❖ Goals:
  - Maximize resource utilization
  - Minimize response time
  - Maximize throughput
- ❖ Some of the popular load balancers are Nginx, ELB (Elastic Load Balancing), HAProxy and Loadbalancer.org.





# What is the N+1 Query Problem?

- ❖ The N+1 query problem happens when a query is executed on each result of the previous query.
- ❖ Here's an example with Rails:
  - Let's say we have `Post` model which belongs to a `User`.

```
Post.all.each do |post|  
  puts "#{post.title} was written by #{post.user.username}"  
end
```

- In the above example, it first retrieved all the `Post` objects and then `user` for each `post`.
- If there are 10 `posts` in the database then  $10 + 1 = 11$  queries would be executed.

# Solve the N+1 problem with "eager loading"

- ❖ With eager loading, a query loads a resource as soon as the code is executed.
- ❖ In Rails, you can use `includes` method for eager loading.
- ❖ Let's rewrite the previous example using `includes` method:

```
Post.includes(:user).each do |post|  
  puts "#{post.title} was written by #{post.user.username}"  
end
```

- ❖ Here, `user` for all the `posts` are pre-loaded in the memory which omits the additional DB calls to fetch user's name.

# Find the N+1 query problem with the Bullet Gem


- ❖ Simple and easy to integrate
- ❖ Automate N+1 tracing
- ❖ Log recorded issues

Tracing example:

```
2009-08-25 20:40:17[INFO] USE eager loading detected:
  Post => [:comments].
  Add to your query: .includes([:comments])
2009-08-25 20:40:17[INFO] Call stack
  /Users/richard/Downloads/test/app/views/posts/index.html.erb:8:in `each'
  /Users/richard/Downloads/test/app/controllers/posts_controller.rb:7:in `index'
```

# Memory Leak / Bloat

- ❖ Memory bloat:
  - Sudden increase in memory consumption
- ❖ Memory leak:
  - More of a slow, gradual increase in memory usage



# Locate Memory Leak / Bloat with APM

- ❖ APM (Application Performance Monitoring) helps to track errors and monitor applications with insights on CPU and memory usage.
- ❖ Some of the popular APM tools are New Relic and Datadog.

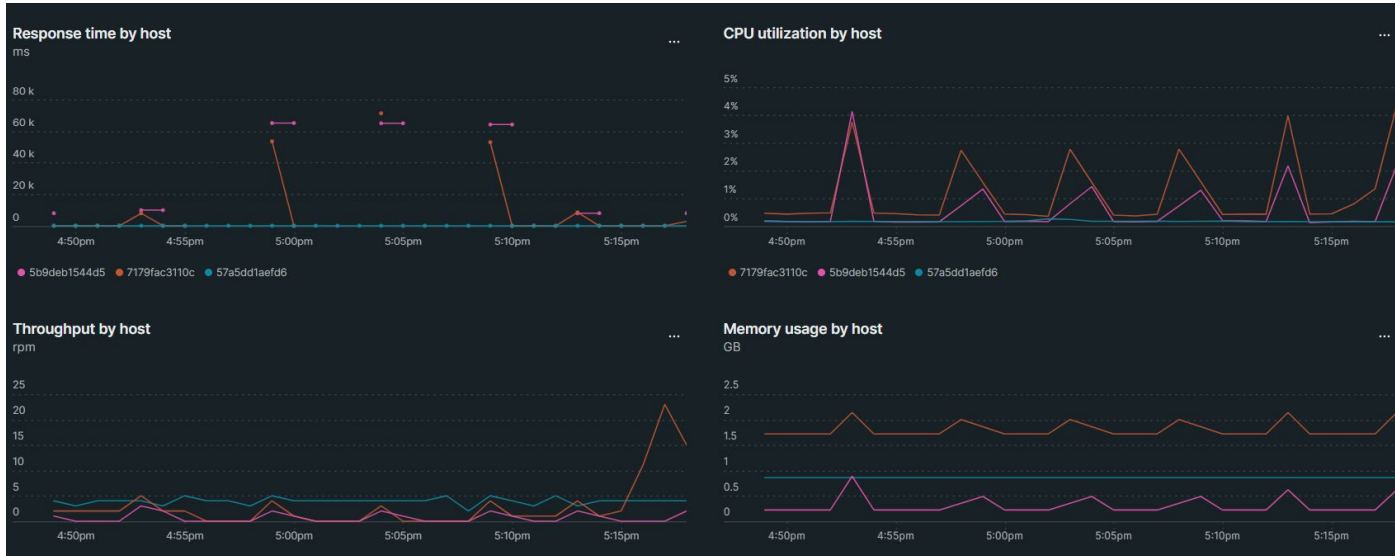



Fig: New Relic Chart

# Profile Rails Endpoints with Skylight.io

- ❖ Besides APMs, Skylight.io is a smart profiler for Ruby and Rails applications which can aid in pinpointing problem endpoints.



Endpoints (44)

Search endpoints...

Name	Response	Typical Response	Problem Response	Popularity	Agony
Api::V1::SportsController#index	jsonapi	98 ms	146 ms		!!
Api::V1::UsersController#show	jsonapi	68 ms	150 ms		!!

Fig: Skylight.io Dashboard

- ❖ Helps locate APIs with:
  - Slow response
  - High object allocations
  - N+1 query

# Cutting down Bulky APIs into Snappy Ones

- ❖ Bulky APIs issues:
  - Slow response cycle
  - Poor experience for end users
  - Heavy server resource usage



- ❖ Snappy APIs benefits:
  - Quick response cycle
  - Good user experience
  - Distributed server load



# Conclusion

- ❖ Facing multiple issues in our experience of scaling a Rails application, we were able to make the following improvements:
  - There was around a 146% increment in the throughput
  - The response time was decreased by 190% in an average
  - Average memory consumption was brought down to ~2.5 GB from ~8 GB
  - Got rid of CPU spikes and server crashes



THANK YOU.

Feel free to ask questions.

# Useful Links

**Premature Scaling:**

<https://bizxpand.com/go-to-market/premature-scaling/#:~:text=What%20exactly%20is%20premature%20scaling,the%20rest%20of%20the%20operation.%E2%80%9D>

**System Scaling:** <https://www.lucidchart.com/blog/what-does-scalability-mean-for-systems-and-services>

**Rails N+1 queries and eager loading:** <https://dev.to/junko911/rails-n-1-queries-and-eager-loading-10eh>

**Load Balancing:** <https://www.nginx.com/resources/glossary/load-balancing/>



Find me on:

- twitter: @as\_m\_it
- linkedIn: <https://www.linkedin.com/in/asmit-ghimire/>